# Malacology: A Programmable Storage System

Paper ID: 165

Total Page: 14

Abstract

Storage systems are caught between the need to evolve data processing applications efficiently and quickly, and the increasing velocity with which storage device technology evolves. This puts tremendous pressure on storage systems to support rapid change both in terms of their interfaces and their performance. But adapting storage systems can be difficult because unprincipled changes might jeopardize years of code-hardening and performance optimization efforts that were necessary for users to entrust their data to the storage system. In this work, we introduce the concept of a programmable storage system, a system that exposes internal services and abstractions with the goal of re-using and building new services on top of the storage stack. We present Malacology, a prototype implementation of a programmable storage system based on Ceph to explore how existing abstractions of common storage system services can be leveraged to adapt to the needs of new data processing systems and the increasing speed of storage devices. We show how this approach allows a large degree of flexibility for storage systems to evolve without sacrificing the robustness of their code-hardened subsystems. We illustrate the advantages and challenges of programmability by composing existing primitives into two new higher-level services: a file system metadata load balancer and a high-performance distributed shared-log that leverages flash devices. The evaluation demonstrates that our services inherit desirable qualities of the back-end storage system, including the ability to balance load, efficiently propagate cluster metadata, recover from failure, and to navigate trade-offs between latency and throughput using leases.

Figure 1: Scalable storage systems have storage daemons which store data, monitor daemons (M) that maintain cluster state, and service-specific daemons (*e.g.*, file system metadata servers). Malacology enables the programmability of internal abstractions (bold arrows) to re-use and compose existing subsystems. With Malacology, we add new services, ZLog and Mantle, that sit alongside traditional user-facing APIs (file, block, object).

1. Introduction

A storage system implements abstractions designed to persistently store data and must exhibit a high level of correctness to prevent data loss. Storage systems have evolved around storage devices that often were orders of magnitude slower than CPU and memory, and therefore could dominate overall performance if not used carefully. Over the last few decades members of the storage systems community have developed clever strategies to meet correctness requirements while somewhat hiding the latency of traditional storage media [12]. To avoid lock-in by a particular vendor, users of storage systems have preferred systems with highly standardized APIs and lowest common denominator abstract data types such as blocks of bytes and byte stream files [3].

A number of recent developments have disrupted traditional storage systems. First, the falling prices of flash storage and the availability of new types of non-volatile memory that are orders of magnitude faster than traditional spinning media are moving overall performance bottlenecks away from storage devices to CPUs and networking, and pressure storage systems to shorten their code paths and incorporate new optimizations [21, 22]. Second, emerging "big data" applications demand interface evolution to support flexible consistency as well as flexible structured data representations. [2]. Finally, production-quality scalable storage systems available as open source software have established and

are continuing to establish new, *de-facto* API standards at a faster pace than traditional standards bodies [26, 32].

The evolutionary pressure placed on storage systems by these trends raises the question of whether there are principles that storage systems designers can follow to evolve storage systems efficiently, without jeopardizing years of code-hardening and performance optimization efforts. In this paper we investigate an approach that focuses on identifying and exposing existing storage system resources, services, and abstractions that in a generalized form can be used to *program* new services. By composing higher-level services over existing abstraction one can reuse subsystems and leverage their optimizations, established correctness, robustness, and efficiency.

We define a programmable storage system to be a storage system that facilitates the re-use and extension of existing storage abstractions provided by the underlying software stack, to enable the creation of new services via composition. A programmable storage system can be realized by exposing existing functionality (such as metadata services and synchronization and monitoring capabilities) as interfaces that can be "glued together" in a variety of ways using a high-level language. In contrast to programmable storage, *active storage* differs through its focus on the injection and execution of code within a storage system or storage device. Programmable storage differs from *active storage* [28]—the injection and execution of code within a storage system or storage device—in that the former is applicable to any component of the storage system, while the latter focuses at the data access level (i.e. inject and execute arbitrary code at the I/O access layers). Given this contrast, we can say that active storage is an example of how one internal component (the storage layer) is exposed in a programmable storage system.

To illustrate the benefits and challenges of this approach we have designed and evaluated Malacology, a programmable storage system that facilitates the construction of new services by re-purposing existing subsystem abstractions of the storage stack. We build Malacology in Ceph, a popular open source software storage stack. We choose Ceph to demonstrate the concept of programmable storage because it offers a broad spectrum of existing services, including distributed locking and caching services provided by metadata servers, durability and object interfaces provided by the backend object store, and propagation of consistent cluster state provided by the monitoring service (see Figure 1). As we will show in this paper, Malacology is expressive enough to provide the functionality necessary for implementing new services.

Malacology includes a non-exhaustive set of interfaces that can be used as building blocks for constructing novel storage abstractions, including:

1. An interface for managing strongly-consistent time-varying **service metadata**.

2. An interface for installing and evolving domain-specific, cluster-wide **data interfaces**.

3. An interface for managing access to **shared resources** using a variety of optimization strategies.

4. An interface for **load balancing** resources across the cluster.

5. An interface for **durability** that persists policies using the underyling storage stack's object store.

We implement two distributed services using Malacology to demonstrate the feasibility of the programmable storage approach:

1. A high-performance distributed shared log service called ZLog, that is an implementation of CORFU. [6]

2. An implementation of Mantle, the programmable load balancing service [30]

The remainder of this paper is structured as follows. First, we describe and motivate the need for programmable storage by describing current practices in the open source community. Next we describe Malacology by presenting the subsystems within the underlying storage system that we re-purpose, and briefly describe how those system are used within Malacology (§4). Then we describe the services that we have constructed within the Malacology framework (§5), and evaluate our ideas within our prototype implementation (§6). We conclude by discussing related and future work.

2. Application-Specific Storage Stacks

Building storage stacks from the ground up for a specific purpose results in the best performance. For example, HDFS was designed specifically to serve Hadoop's jobs, and uses techniques like exposing data locality and relaxing POSIX constraints to to achieve application-specific I/O optimizations. [**?**]. Alternatively, general-purpose storage stacks are built with the flexiblity to serve many applications by providing a variety of interfaces and tunable parameters. Unfortunately, managing competing forces in these systems is difficult and users want more control from the general-purpose storage stacks without going as far as building their storage system from the ground up.

To demonstrate this trend in storage systems we examine the state of programmability in Ceph. Something of a storage swiss army knife, Ceph simultaneously supports file, block, and object interfaces on a single cluster [15]. Ceph's Reliable Autonomous Distributed Object Storage (RADOS) system is a cluster of storage devices (OSDs) that provide Ceph with data durability and integrity using replication, erasure-coding, and scrubbing [40]. Ceph already provides some degree of programmability; the OSDs support domain-specific object interfaces that are implemented by composing existing low-level storage interfaces that execute atomi-

cally. These interfaces are written in C++ and are statically loaded into the system.

The Ceph community provides empirical evidence that developers are already beginning to embrace programmable storage. Figure 2 shows a dramatic growth in the production use of domain-specific interfaces in the Ceph community since 2010. In that figure, classes are functional groupings of methods on storage objects (e.g. remotely computing and caching the checksum of an object extent). What is most remarkable is that this trend contradicts the notion that API changes are a burden for users. Rather it appears that a gap in existing interfaces are being addressed through ad-hoc approaches to programmability. In fact, Table 1 categorizes existing interfaces and we clearly see a trend towards reusable services.

Figure 2: [source] Since 2010, the growth in the number of co-designed object storage interfaces in Ceph has been accelerating. This plot is the number of object classes (a group of interfaces), and the total number of methods (the actual API end-points).

Category	Example	#
Logging	Geographically distribute replicas	11
Metadata Management	Snapshots in the block device OR Scan extents for file system repair	74
Locking	Grants clients exclusive access	6
Other	Garbage collection, reference counting	4

Table 1: A variety of object storage classes exist to expose interfaces to applications. # is the number of methods that use these categories.

The takeaway from Figure 2 is that programmers are already trying to use programmability because their needs, whether they be related to performance, availability, consistency, convenience, etc., are not satisfied by the existing default set of interfaces. The popularity of the custom object interface facility of Ceph could be due to a number of reasons, such as the default algorithms/tunables of the storage system being insufficient for the application's performance goals, programmers wanting to exploit application-specific semantics, and/or programmers knowing how to manage resources

to improve performance. A solution based on application-specific object interfaces is a way to work around the traditionally rigid storage APIs because custom object interfaces give programmers the ability to tell the storage system about their application: if it is CPU or IO bound, if it has locality, if its size has the potential to overload a single proxy node, etc. The programmers know what the problem is and how to solve it, but until the ability to modify object interfaces, they had no way to express to the storage system how to handle their data.

Our approach is to expose more of the commonly used, code-hardened subsystems of the underlying storage system as interfaces. The intent is that these interfaces, which can be as simple as a redirection to the persistent data store or as complicated as a strongly consistent directory service, should be used and re-used in many contexts to implement a wide range of services. By making programmability a 'feature', rather than a 'hack' or 'workaround', we help standardize a development process that now is largely ad-hoc.

3.   Challenges

Implementing the infrastructure for programmability into existing services and abstractions of distributed storage systems is challenging, even if one assumes that the source code of the storage system and the necessary expertise for understanding it is available. Some challenges include:

- Storage systems are generally required to be highly available so that any complete restarts of the storage system to reprogram them is usually unacceptable.

- Policies and optimizations are usually hard-wired into the services and one has to be careful when factoring them out not to introduce additional bugs. These policies and optimiations are usually cross-cutting solutions to concerns or trade-offs that cannot be fully explored yet (as they relate to workload or hardware). Given these policies and optimizations, decomposition of otherwise orthogonal internal abstractions can be difficult or dangerous.

- Mechanisms that are often only exercised according to hard-wired policies and not in their full generality have hidden bugs that are revealed as soon as those mechanisms are governed by different policies. In our experience introducing programmability into a storage system proved to be a great debugging tool.

- Programmability, especially in live systems, implies changes that need to be carefully managed by the system itself, including versioning and propagation of those changes without affecting correctness.

To address these challenges we present Malacology. Malacology is both a prototype for a programmable storage system and a design approach to evolve storage systems efficiently and without jeopardizing years of code-hardening

and performance optimization efforts. Although Malacology uses the internal abstractions of the underlying storage system, including its subsystems, components, and implementations, we emphasize that our system still addresses the general challenges outlined above.

The main challenge of designing a programmable storage system is choosing the right internal abstractions and picking the correct layers for exposing them. In this paper, we do not present an exhaustive list of possible internal abstractions nor do we contend that the abstractions we choose provide the best trade-offs for all applications. For example, if consensus is correctly exposed one could implement high-level features like versioning, serialization, or various flavors of strongly consistent data management on top; but perhaps a low-level consensus interface is suited well for a particular set of applications. These questions are not answered in this paper and instead we focus on showing the feasibility of building such a system, given advances in the quality and robustness of today's storage stacks.

Malacology implements and demonstrates the ideas of programmable storage in Ceph. We choose Ceph because it is a production quality system and because it is open source. The large developer community ensures that code is robust and the visibility of the code lets us expose any interface we want. In the next section we describe the Ceph components that we expose as Malacology interfaces.

4. Malacology: A Programmable Storage System

The guiding principle is to re-use existing services and extend them so that these services can be *programmed*. We accomplish programmability of a service by exporting bindings (or "hooks") for an interpreted programming language so that programming can occur without having to restart the storage system (see also below, §4.4).

There are multiple reasons that make Lua an attractive runtime for implementing Malacology. Lua is a portable, embedded scripting language that offers superior performance and productivity trade-offs, including a JIT-based implementation that is well known for near native performance. Additionally, Lua has been used extensively in game engines, and systems research [36], including storage systems where it has been effectively used both on [18, 20, 39] and off [30] the performance critical path. Finally, the flexibility of the runtime allows execution sandboxing in order to address security and performance concerns.

We will now discuss the common subsystems used to manage storage system and how Malacology makes them programmable.

4.1 Service Metadata Interface

Keeping track of the state of a distributed system is an essential part of any successful service and a necessary component in order to diagnose and detect failures, when they occur. This is further complicated by variable propagation delays and heterogenous hardware in highly dynamic environments.

In the case of Ceph, a consistent view of cluster state among server daemons and clients is critical to provide strong consistency guarantees to clients. Ceph maintains cluster state information in per-subsystem data structures called "maps" that record membership and status information. A Paxos-based monitoring service is responsible for integrating state changes into cluster maps, responding to requests from out-of-date clients and synchronizing members of the cluster whenever there is a change in a map so that they all observe the same system state. As a fundamental building block of many system designs, consensus abstractions such as Paxos are a common technique for maintaining consistent data versions, and are a useful system to expose.

The default behavior of the monitor can be seen as a Paxos-based notification system, similar to the one introduced in [13], allowing clients to identify when new values (termed epochs in Ceph) are associated to given maps. While Ceph does not expose this service directly, a key-value service designed for managing configuration metadata that is built on top of the consensus engine is an example of a high-level service within Ceph that re-uses existing abstractions. While a key-value service is generally useful, it doesn't provide many of the useful services hidden within the monitoring framework that would be required for applications with more demanding requirements, such as creating arbitrary maps with application-specific semantics and interfaces. Since the monitor is intended to be out of high-performance I/O paths, a general guideline is to make use of this functionality infrequently and to assign small values to maps.

Malacology: Malacology exposes a strongly-consistent view of time-varying service metadata as a service rather than a hidden internal component. Malacology provides a generic API for adding arbitrary values to existing subsystem cluster maps. As a consequence of this, applications can define simple but useful service-specific logic to the strongly-consistent interface, such as authorization control (just specific clients can write new values) or to trigger actions based on specific values (e.g. sanitize values). The higher level services we implement in §5 make use of this functionality to register, version and propagate dynamic code (Lua scripts) for new object interfaces defined in storage daemons (§4.2) and policies in the load balancer §4.3). Using this service guarantees that interface definitions are not only made durable, but are transparently and consistently propagated throughout the cluster so that clients are properly synchronized with the latest interfaces.

4.2 Object I/O Interfaces

Briefly described in Section 2, Ceph supports the creation of application-specific object interfaces [40]. The ability to

Common Internal Abstractions

Abstraction	Example in Production Systems	Example in Ceph	Provided Primitives
Service Metadata	Zookeeper/Chubby coordination [11, 24]	cluster state management [34]	consensus & consistency
Shared Resource	MPI collective IO, burst	POSIX metadata protocols	serialization & batching
Object & Data	Swift in situ storage/compute [1]	object interface classes [37]	transaction & atomicity
Durability	S3/Swift interfaces (RESTful API)	object store library [40]	interface to persistence
Load Balancing	VMWare's VM migration [17, 23]	migrate POSIX metadata [41]	migration & sampling

Table 2: The same "internal abstractions" are common across large-scale systems because they provide primitives that solve general distributed systems problems. Here we list examples of what these internal abstractions are used for in "production systems" and in Ceph. Malacology provides these internal abstractions as interfaces that higher level applications can use; *italicized abstractions* are interfaces contributed in this paper.

offload computation can reduce data movement, and transactional interfaces can significantly simplify construction of complex storage interfaces that require uncoordinated parallel access.

An object interface is a plugin structured in a similar way to that of an RPC in which a developer creates a named function within the cluster that clients may invoke. In the case of Ceph each function is implicitly run within the context of an object specified when the function is called. Developers of object interfaces express behavior by creating a composition of native interfaces or other custom object interfaces, and handle serialization of function input and output. A wide range of native interfaces are available to developers such as reading and write to a byte stream, controlling object snapshots and clones, and accessing a sorted key-value database. These native object interfaces may be transactionally composed along with application specific logic to create interfaces with a rich set of semantics. An example is the use of a read-modify-write operation to perform consistent object updates across the byte stream and key-value database interfaces without relying on complex client-driven concurrency control.
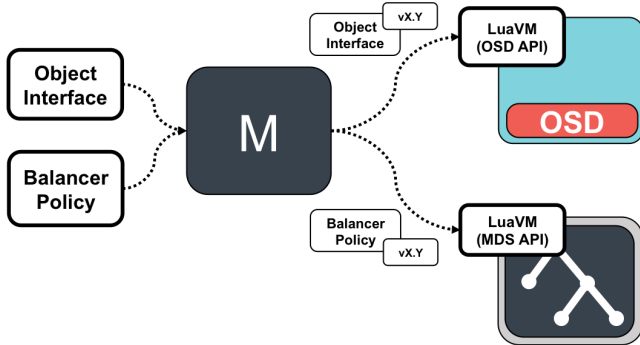
Figure 3: Malacology re-uses Ceph subsystems to allow users to dynamically define new object/data interfaces and load balancing policies. It achieves this by extending the object storage daemon (OSD) and metadata server daemon (MDS) subsystems with an embedded Lua VM. It uses the of the strongly-consistent service metadata interface to propogate the object interfaces and versions across the cluster.

Malacology: The implementation of Ceph's object abstraction, although powerful, does not readily support programmability. Supporting only C/C++ for object interface developers, Ceph requires distribution of compiled binaries for the correct architecture, adding a large barrier of entry for developers and system administrators. Second, having no way to dynamically unload modules, any changes require a full restart of a storage daemon which may have serious performance impacts due to loss of cached data. And finally, the security limitations of the framework limit the use of object interfaces to all but those with administrative level access and deep technical expertise.

To address these concerns, Malacology takes advantage of Lua extensions contributed by the Ceph community. This allows new object interfaces to be dynamically loaded into the system and modified at runtime, resulting in a object storage API with economy of expression, which at the same time provides the full set of features of the base object class. New object interfaces that are expressed in thousands of lines of code can be implemented in approximately an order of magnitude less code [18]. While the use of Lua does not prevent deployment of malicious code, certain types of coding mistakes can handled gracefully, and access policies are used to restrict the feature to trusted users.

All of these modifications work in tandem to implement the desired behavior, and are typically co-designed together such that each depends on the other to behave as expected.

4.3 Distributed Metadata Interfaces

The distributed metadata service in Ceph provides clients with a POSIX file system abstraction [41]. In general, distributed file systems protect resources by providing hierarchical indexing and distributed locking services. In Ceph, the locking service implements a capability-based system that expresses what data and metadata clients are allowed to access as well as what state they may cache and modify locally. While designed for a fixed file abstraction, indexing, locking, and caching are all common services that are useful to a broad spectrum of applications. Distributed applications that share centralized resources (e.g. a database

or directory) face similar challenges which are often solved using application-specific sharding.

Ceph also addresses the challenge of balancing metadata load with a separate metadata cluster. This cluster uses load balancing policies to migrate directory inodes around the cluster to alleviate overload at single nodes. The policies use metrics based on system state (e.g. CPU and memory utilization) and statistics collected by the cluster (e.g. the popularity of an inode). Ceph uses dynamic subtree partitioning to move variable sized namespace subtrees [41]. These units can be shipped anywhere (i.e., to any metadata server of any capacity) at any time for any reason. The original balancer was designed with hard-coded policies and tunables.

Malacology: interfaces are added in strategic spots in the metadata service for guarding shared resources, defining new file types, and load balancing.

Shared resource interface. The metadata service manages client sessions, allowing clients to obtain locks (e.g. file byte ranges), and capabilities (e.g. to cache file data). Clients and metadata servers use a cooperative protocol in which clients voluntarily release resources back to the metadata service in order to implement sharing policies. While the current policy for sharing access and voluntarily releasing resources is largely best-effort, Malacology supports generalized policies between metadata servers and clients that can be used to implement fairness or priority.

File type interface. The metadata service manages a POSIX file system hierarchy where files and directories are represented as inode data structures and expose a POSIX file interface. Application that manage large amounts of metadata (e.g. users or database snapshots) often require a naming service. We allow new inode types to be defined such that applications can create domain-specific interfaces to inodes that may modify locking and capability policies. We will show this is used in Section 5.2.1 when we discuss a distributed shared-log built on Malacology.

Load balancing interface. Mantle [30], a programmable metadata load balancer, has been implemented on Malacology so that its infrastructure can enjoy all the properties of existing internal abstractions. The load balancing logic is intact but the infrastructure has been improved to safely load, manage, and persist its load balancing policies. Composing Mantle using the service metadata and durability interfaces, Mantle can safely version balancer policies, save balancer policies in the backend object store and centralize warnings/errors. When combined with file types, load balancing policies can express a variety of rules such as multi-tentant and multi-application balancing.

4.4 Durability Interface

Ceph provides storage by striping and replicating data across RADOS [40], the reliable distributed object store. RADOS uses many techniques to ensure that data is not corrupted or lost, such as erasure coding, replication, and data scrubbing.

Service	Used to...	Malacology Interface
Mantle	load balance metadata	Load Balancing
	agree on current balancer	Service Metadata
	control client write lock	Shared Resource
	persist policy logic	Durability
ZLog	implement CORFU	Object & Data
	install ZLog interfaces	Service Metadata
	persist log entries	Durability
ZLog Seqr.	load balance sequencers	Load Balancing
	manage access to log tail	Shared Resource

Table 3: Mantle and ZLog are built using Malacology interfaces so they inherit the robustness of the subsystems of the underlying stack.

Furthermore, many of these features are implemented using a peer-to-peer protocol that allows object storage devices to operate autonomously without a centralized coordinator. For example, when the number of placement groups change, the OSDs re-balance and re-shard data in the background in a process called placement group splitting in which OSDs communicate directly with each other to coverage on a new state. In order to reduce load on the monitoring service, Ceph OSDs use a gossip protocol to efficiently propagate changes to cluster maps throughout the system, and autonomously initiate recovery mechanisms when failures are discovered.

Malacology: Metadata service policies and object storage interfaces are stored durability within RADOS and managed by storing references with the object maps. Since the cluster already propagates a consistent view of these data structures, we use this service to automatically install interfaces in OSDs, and install policies within the metadata server daemon (MDS) such that clients and daemons are synchronized on correct implementations without restarting.

5. Services Built on Malacology

In this section we describe two services built on top of Malacology. The first is Mantle, a framework for dynamically specifying metadata load balancing policies. The second system, ZLog, is a high-performance distributed shared-log. In addition to these services, we'll demonstrate how we combine ZLog and Mantle to implement service-aware metadata load balancing policies.

5.1 Mantle: Programmable Load Balancer

Mantle is a programmable load balancer that separates the metadata balancing policies from their mechanisms. Administrators inject code to change how the metadata cluster distributes metadata. In [30] the authors showed how using Mantle, one can implement a single node metadata service, a distributed metadata service with hashing, and a distributed metadata service with dynamic subtree partitioning.

The original implementation was "hard-coded" into Ceph and lacked robustness (no versioning, durability, or policy distribution). Re-implemented using Malacology, Mantle

now enjoys (1) the versioning provided by the metadata service interface provided by the monitor daemons and (2) the durability and distribution provided by Ceph's reliable object store.

5.1.1 Versioning Balancer Policies

Ensuring that the version of the current load balancer is consistent across the physical servers in the metadata cluster was ignored in the original implementation. The user had to set the version on each individual server and it was trivial to make the version inconsistent. Maintaining a consistent version is important for balancing policies that are cooperative in which local decisions are made assuming properties about other instances of the same policy.

With Malacology, Mantle stores the version of the current load balancer in the Ceph service metadata interface. The version of the load balancer corresponds to an object name in the balancing policy. Using the service metadata interface ensures that all metadata servers use the same version of the load balancer. As a result, the policy version inherits the consistency of Ceph's internal monitor daemons.

The user changes the version of the load balancer using a new CLI command:

```
$> ceph fs set lua_balancer_class <version>
```

5.1.2 Making Balancer Policies Durable

The load balancer version described above corresponds to the name of an object in RADOS that holds the actual Lua balancing code. When MDS nodes start balancing load, they first check the latest version from the MDS map and compare it to the balancer they have loaded. If the version has changed, they dereference the pointer to the balancer version by reading the corresponding object in RADOS. This is in contrast to the original balancer which stored load balancer code on the local file system – a technique which is unreliable and may result in silent corruption.

The balancer pulls the Lua code from RADOS synchronously; asynchronous reads are not possible because of the architecture of the MDS. The synchronous behavior is not the default behavior for RADOS operations, so we achieve this with a timeout: if the asynchronous read does not come back within half the balancing tick interval the operation is cancelled and a Connection Timeout error is returned. By default, the balancing tick interval is 10 seconds, so Mantle will use a 5 second second timeout.

This design allows Mantle to immediately return an error if anything RADOS-related goes wrong. We use this implementation because we do not want to do a blocking OSD read from inside the global MDS lock. Doing so would bring down the MDS cluster if any of the OSDs are not responsive.

Storing the balancers in RADOS is only possible because we use Lua as the language for writing balancer code. If we used a language that needs to be compiled, like the C++ object classes in the OSD, we would need to ensure binary compatibility, which is complicated by different operating systems, distributions, and compilers.

5.1.3 Logging, Debugging, and Warnings

In the original implementation, Mantle would log all errors, warnings, and debug messages to a log stored locally on each MDS server. To get the simplest status messages or to debug problems, the user would have to log into each MDS individually, look at the logs, and reason about causality and ordering.

With Malacology, Mantle re-uses the centralized logging features of the monitoring service. Important errors, warnings, and info messages are collected by the monitoring subsystem and appear in the monitor cluster log so instead of users going to each node, they can watch messages appear at the monitor. Messages are logged sparingly, so as not to spam the monitor with frivolous debugging but important events, like balancer version changes or failed subsystems show up in the centralized log.

5.2 ZLog: A Fast Distributed Shared Log

The second service implemented on Malacology is ZLog, a high-performance distributed shared-log that is based on the CORFU protocol [6]. A shared-log is a general, yet powerful abstraction used to construct distributed systems, such as cloud-based metadata management [5] and elastic database storage engines [7–9]. However, existing implementations that rely on consensus algorithms such as Paxos funnel I/O through a single point introducing a bottleneck that restricts throughput. In contrast, the CORFU protocol is able to achieve high throughput using a network counter service, called a *sequencer*, that is used to decouple log position assignment from log I/O.

While a full description of the CORFU system is beyond the scope of this paper, we briefly describe the custom storage device interface, sequencer service, and recovery protocol, and how these services are instantiated in the Malacology framework.

5.2.1 Sequencer

High-performance in CORFU is achieved using a sequencer service that assigns log positions to clients by reading from a volatile, in-memory counter which can run at a very high throughput and at low latency. Since the sequencer is centralized, ensuring serializability in the common case is trivial. The primary challenge in CORFU is handling the failure of the sequencer in a way that preserves correctness. Failure of the sequencer service in CORFU is handled by a recovery algorithm that recomputes the new sequencer state using an application-specific custom storage interface to discover the tail of the log, while simultaneously invalidating stale client requests using an epoch-based protocol.

Sequencer interface. The sequencer resource supports the ability to *read()* the current tail value as well as get the *next()* position in the log which also atomically incre-

ments the tail position. We implement the sequencer service in Malacology as an application-specific file type that associates a small amount of metadata, including the 64-bit log tail position, to the inode state managed by the Ceph metadata service. Since a sequencer is associated with a unique log, this approach has the added benefit of allowing the metadata service to also handle naming, by representing each log service instance within the standard POSIX hierarchical namespace. The primary challenge in mapping the sequencer resource to the metadata service is handling serialization correctly to maintain the global ordering provided by the CORFU protocol.

Initially we sought to directly model the sequencer service in Ceph as a non-exclusive, non-cacheable resource, forcing clients to perform a round-trip access to the resource at the authoritative metadata server for the sequencer inode. Interestingly, we found that the capability system in Ceph uses a strategy to reduce metadata service load by allowing clients that open a shared file to temporarily obtain an exclusive cached copy of the resource, resulting in a round-robin, best-effort batching behavior. When a single client is accessing the sequencer resource it is able to increment the sequencer locally. Any competing client cannot query the sequencer until the metadata service has granted it exclusive access.

While unexpected, this discovery allowed us to explore an implementation strategy that we had not previously considered. In particular, for bursty clients, and clients that can tolerate increased latency, this mode of operation may allow a system to achieve much higher throughput than a system with a centralized sequencer service. We utilize the programmability of the metadata service to define a new policy for handling capabilities that controls the amount of time that clients are able to cache the sequencer resource. This allows an administrator or application to control the trade-off between latency and throughput beyond the standard best-effort policy that is present in Ceph by default.

In Section **??** we quantify the trade-offs of throughput and latency for an approach based on a round-robin batching mode, and compare this mode to one in which the metadata server mediates access to the sequencer state when in it is being shared among multiple clients. Quantifying these trade-offs should provide administrators with guidelines for setting the tunables for different "caching" modes of the sequencer.

Balancing policies. As opposed to the batching mode for controlling access to the sequencer resource, more predictable latency can be achieved by treating the sequencer inode as a shared non-cacheable resource, forcing clients to make a round-trip to the metadata service. However, the shared nature of the metadata service may prevent the sequencer from achieving maximum throughput. To provide a solution to this issue we have taken advantage of the programmability of the metadata service load balancing to con-

struct a service-specific load balancing policy. As opposed to a generic balancing policy that may strive for uniform load distribution, a ZLog-specific policy may utilize knowledge of inode types to migrate the sequencer service to provisioned hardware during periods of contention or high demand.

5.2.2 Storage Interface

The storage interface is a critical component in the CORFU protocol. Clients independently map log positions that they have obtained from the sequencer service (described in detail in the next section) onto storage devices, while storage devices provide an intelligent *write*-once, random *read* interface for accessing log entries. The key to correctness in CORFU lies with the enforcement of up-to-date epoch tags on client requests; requests tagged with out-of-date epoch values are rejected, and clients are expected to request a new tail from the sequencer after refreshing state from an auxiliary service. This mechanism forms the basis for sequencer recovery.

In order to repopulate the sequencer state (i.e. the cached, current tail of the log) during recovery of a sequencer, the maximum position in the log must be obtained. To do this, the storage interface exposes an additional *seal* method that atomically installs a new epoch value and returns the maximum log position that has been written.

Since the sequencer servivce does not resume until the recovery process has completed, there cannot be a race with clients appending to the log, and the immutability of the log allow reads to never block during a sequencer failure. Recovery of the sequencer process itself may be handled in many ways, such as leader election using an auxiliary service such as Paxos. In our implementation, the recovery is the same (and is inherited from) the Ceph metadata service. Handling the failure of a client that holds the sequencer state is similar, although a timeout is used to determine when a client should be considered unavailable.

6. Evaluation

Our evaluation demonstrates the feasibility of building new service abstractions atop programmable storage, focusing on the performance of the internal abstractions exposed by Malacology and used to construct the Mantle and ZLog services. We also discuss latent capabilities we discovered in this process that let us navigate different trade-offs within the services themselves. First, we benchmark scenarios with high sequencer contention by examining the interfaces used to map ZLog onto Malacology; specifically, we describe the sequencer implementation and the propagation of object and data interfaces interfaces. Next, we benchmark scenarios in which the storage system manages multiple logs by using Mantle to balance sequencers across a cluster.

Since this work focuses on the programmabilty of Malacology, the goal of this section is to show that the compo-

nents and subsystems that support the Malacology interfaces provide reasonable relative performance, as well as to give examples of the flexibility that Malacology provides to programmers.

6.1 Mapping ZLog onto Malacology

We evaluate Malacology by exploring one possible mapping of the ZLog implementation of CORFU onto Ceph in which we re-use (1) the metadata service to manage naming and synchronization of the sequencer resource by treating the resource as an inode, and (2) the monitoring sub-subsystem to distribute and install application-specific I/O interface required of the CORFU protcool. In Section 6.2 we then demonstrate how the re-use of the inode abstraction for implementing the sequencer resource enables load balancing policies to migrate the sequencer resource in heavy-load situation.

6.1.1 Sequencer Implementation

We evaluate the feasibility of using the metadata service to implement a seqeuencer resource that is responsible for maintaining a total ordering of the log. Clients contact the sequencer to obtain the tail of the log and then independently initiate I/O, thus we measure both the throughput and latency of obtaining new tail positions which inherently bound client append performance.

The sequencer is implemented as a custom file type in which the sequencer state (a 64-bit integer) is embedded in the inode of a file representing the sequencer. A total ordering of the log is imposed by the re-use of the capability service that can be used to grant exclusive access of inode state to clients. The metadata service is responsible for maintaining exclusivity and granting access. Figure 4 (a) shows the behavior of the system in which a best-effort policy is used. The two colors represent points in time that the clients were able to access the resource. The best-effort policy provides interleaving a high degree of interleaving between clients but the system spends a large portion of time re-distributing the capability, reducing overall throughput.

In order to control the performance of the system we implement a policy that (1) restricts the length of time that a client may maintain exclusive access and (2) limits the number of log positions that a client may generate without yielding to other clients waiting for access. The behavior of these two modes is illustrated in Figures 4 (b) and (c), respectively.

Figure 5 demonstrates a configurable trade-off between throughput and latency. In the experiment two clients are run each with a fixed 0.25 second maximum reservation on the capability, and we vary the size of the log position quota running each configuration for two minutes. The total operations per second is the combined throughput of the two clients, and the average latency is number of microseconds required to obtain a new log position. With a small quota more time is spent exchanging exclusive access, while a
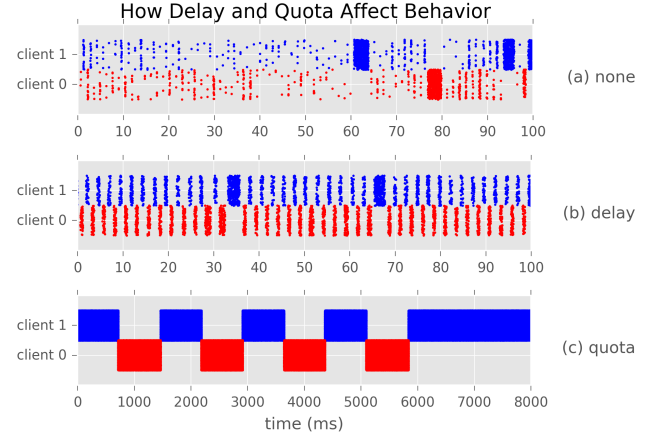
Figure 4: Each dot in these graphs is an individual request and they are spread randomly along the y axis. The default behavior ("None") shows unpredictable performance, "delay" lets each client hold the lease longer, and "quota" gives each client the lease for a certain number of operations. Delay improves latency and fairness while quota maximizes throughput.

large quota reservation allows clients to experience a much lower latency because they uncontended access for a longer period of time.

Figure 5: [source] Sequencer throughput by re-using various services. The highest performance is achieved using a single client with exclusive, cacheable privilege. Round-robin sharing of the sequencer resource is affected by the amount of time the resource is held, with best-effort performing the worst.

To get a better picture of latency, Figure 6 shows the CDF of latency for each client in all experiment configurations. At the 99th percentile clients accessed the sequencer is less than a millisecond. The CDF is cropped at the 99.999th percentile due to large outliers that we believe occur in instances in which the MDS is performing I/O while it is in the process of re-distributing the capability to another client.

We have shown how exposing the internal capability management service allows users to navigate latency and throughput tradeoffs. Other approaches to designing the sequencer service also exist, such as using a centralized service in which each access is a network round-trip. In contrast to the mechanism we explored which is appropriate for clients with bursty workloads, it may be easier to provide predictable performance using a centralized service, and we will be exploring in the future how this can be achieved using the capability system.

Figure 6: The latency distribution shows that if the clients hold their capabilities longer, performance decreases. Changing the delay via policy can help the sequencer and clients strike a balance between throughput and latency.

Figure 7: Cluster-wide interface update latency, excluding the Paxos proposal cost for interface commit.

6.1.2 Interface Propagation

Domain-specific data interfaces (Section **??**) allow co-design between applications and the storage system. Malacology supports custom object interfaces in RADOS that require interface implementations to be installed on the storage devices in the system, supporting the evolution of interfaces through automatic system-wide versioning and installation through the service metdata interface (Section 4.1). We evaluate the performance of installing a new version of an interface in the cluster, which is an important metric for applications that are designed with frequently evolving interfaces.

We demonstrate the feasibility of utilizing the Ceph monitoring sub-system by evaluating the performance of installing and distributing interface updates. Figure 7 shows the CDF of the latency of interface updates. The interfaces are Lua scripts embedded in the cluster map and distributed using a peer-to-peer gossip protocol. The latency is defined as the elapsed time following the Paxos proposal for an interface update until each OSD makes the update live (the cost of the Paxos proposal is configurable and is discussed below). The latency measurements were taken on the nodes running Ceph server daemons, and thus exclude the client round-trip cost. In each of the experiments 1000 interface updates were observed.

Figure 7 shows the lower bound cost for updates in a large cluster. In the experiment labeled "120 OSD (RAM)" a cluster of 120 OSDs (10 OSDs x 12 servers) using an in-memory data store were deployed, showing a latency of less than 54 ms with a probability of 90% and a worst case latency of 194 ms. These costs demonstrate the penality of distributing the interface in a large cluster. In practice the costs include, in addition to cluster-wide propagation of interface updates, the network round-trip to the interface management service, the Paxos commit protocol itself, and other factors such as system load. By default Paxos proposals occur periodically with a 1 second interval in order to accumulate updates. In a minimum, realistic quorum of 3 monitors using HDD-based storage, we were able to decrease this interval to an average of 222 ms.


Figure 8: CephFS/Mantle load balancing have better throughput than co-locating all sequencers on the same server. Sections 6.2.1 and 6.2.2 quantify this improvement; Section 6.2.3 examines the migration at 0-60 seconds.

6.2 Load Balancing ZLog Sequencers with Mantle

In practice, a storage system implementing CORFU will support a multiplicity of independent totally-ordered logs for each application. For this scenario co-locating sequencers on the same physical node is not ideal. Building a load balancer that can migrate the shared resource (in this case, the resource that mediates access to the tail of the log) is a time-consuming, non-trivial task that requires building subsystems for migrating resources, monitoring the workloads, collecting metrics that describe the utilization on the physical nodes, partitioning resources, maintaining cache coherence, and managing multiple sequencers. The following experiments demonstrate the feasibility of using the mechanisms of the Malacology load balancing interface to inherit all these features and to alleviate load from overloaded servers housing multiple sequencers.

The experiments are run on a cluster with 10 nodes to store objects, one node to monitor the cluster, and 3 nodes that can accomodate sequencers. Our first experiment co-locates 3 sequencers each with 4 clients on the same physical server. Instead of measuring contention at the clients like Section 6.1.1, these experiments measure contention at the sequencers by forcing clients to make round-trips for every request. We implement this using the shared resource interface that forces round-trips. Because the sequencer's only function is to hand out positions for the tail of the log, the workload is read-heavy.

First, we show how the ZLog service can orchestrate multiple sequencers in a cluster using the Malacology load bal-

Figure 9: In (a) all CephFS balancing modes have the same performance; Mantle uses a balancer designed for sequencers. In (b) the best combination of mode and migration units can have up to a $2\times$ improvement.


Figure 10: In client mode clients sending requests to the server that houses their sequencer. In proxy mode clients continue sending their requests to the first server.

ancing interface. Figure 8 shows the throughput (y axis) over time (x axis) of different load balancers that migrate ZLog sequencers acround the cluster; "No Balancing" keeps all sequencers on one server, "CephFS" migrates sequencers using the CephFS load balancers, and "Mantle" uses a custom load balancer we wrote specifically for sequencers. The CephFS load balancers are hard-coded into Ceph while Mantle is an API for balancing load presented in Mantle [30] and re-implemented on top of Malacology in this paper. The increased throughput for the CephFS and Mantle curves between 0 and 60 seconds are a result of migrating the sequencer(s) off overloadeed servers.

In addition to showing that migrating sequencers improves performance, Figure 8 also demonstrates features that we will explore in the rest of this section. Sections 6.2.1 and 6.2.2 quantify the differences in performance when the cluster stabilizes at time 100 seconds and Section 6.2.3 examines the slope and start time of the rebalancing phase between 0 and 60 seconds by comparing the aggressiveness of the balancers.

6.2.1 Feature: Balancing Modes

Next, we quantify the performance benefits shown in Figure 8. To understand why load balancers perform differently we need to explain the different balancing modes that the load balancer service uses and how they stress the subsystems that receive and forward client requests in different ways. In Figure 8, the CephFS curve shows the performance of the balancing mode that CephFS falls into *most of the time*. CephFS has 3 modes for balancing load: CPU mode, workload mode, and hybrid mode. All three have the same structure for making migration decisions but vary based on the metric used to calculate load. For this sequencer workload the 3 different modes all have the same performance, shown in Figure 9, because the load balancer falls into the same mode a majority of the time. The high variation in performance for the CephFS CPU Mode bar reflects the un-

certainty of using something as dynamic and unpredictable as CPU utilization to make migration decisions. In addition to the suboptimal performance and unpredictability, the fact that the CephFS balancers all fall into the same mode despite using different balancers is a problem becuase it prevents administrators from properly exploring the balancing state space.

Mantle gives the administrator more control over the balancing policies; for the Mantle bar in Figure 9 we use the load balancing interface to program logic for balancing read-heavy workloads, resulting in higher throughput and better stability. When we did this we also identified two balancing modes that are relevant for making migration decisions for sequencers.

Using Mantle, the administrator can put the load balancing service into "proxy mode" or "client mode". In proxy mode one server receives all requests and farms off the requests to slave servers; the slaves servers do the actual tail finding operation. In client mode, clients interact directly with the server that has their sequencer. These modes are illustrated in Figure 10: "No balancing" is when all sequencers are co-located on one physical server – performance for that mode is shown by the No Balancing curve in Figure 8. "Client Mode" in the lower left shows how clients for sequencer 2 are re-directed to server B. The two "Proxy mode" diagrams on the right show how clients continue sending reqeusts to server A even though sequencer 2 has migrated to server B (the meaning of "(Half)" is discussed in the next section).

Figure 11 shows the throughput over time of the two different modes for an environment with only 2 sequencers (again 4 clients each) and 2 servers. The curves for both sequencers in Figure 11(a) start at less than 1000 ops/second and at time 60 seconds Mantle migrates Sequencer 1 to the slave server. Performance of Sequencer 2 decreases because it stayed on the proxy which now handles all client requests, processes requests for Sequencer 2, and forwards requests for Sequencer 1. The performance of Sequencer 1 improves dramatically because distributing the sequencers in this way separates (1) the handling of the client requests and (2) finding the tail of the log and responding to clients. Doing both steps is too heavy weight for one server and sequencers on slave nodes can go faster because work is split up; this

Figure 11: The performance of proxy mode achieves the highest throughput but at the cost of lower throughput for one of the sequencers. Client mode is more fair but results in lower cluster throughput.

phenomenon is not uncommon and has been observed in chain replication [**?**]. Cluster throughput improves at the cost of decreased throughput for Sequencer 2. Figure 11(b) is set to sequencer mode manually (no balancing phase) and shows that the cluster throughput is worse than the cluster throughput of proxy mode. That graph also shows that Sequencer 2 has less throughput than Sequencer 2. In this case, the scatter-gather process used for cache coherence in the metadata protocols causes strain on the server housing Sequencer 2 resulting in this uneven performance.

6.2.2 Feature: Migration Units

Another factor that affects performance in this environment is "how much" load is on each server; these experiments quantify that effect by programming the load balancing interface to control the amount of load to migrate. We call this metric a "migration unit". Expressing this heuristic is not easily achievable with outward facing tunable parameters (i.e. system knobs) but with Mantle's programmable interface, we can force the load balancer to change its migration units. To force the balancer into the Proxy Mode (Half) scenario in Figure 10, which uses migration units that are half the load on the current server:

```
targets[whoami+1] = mds[whoami]["load"]/2
```

This code snippet uses globally defined variables and tables from the Mantle API to send half of the load on the current server (whoami) to the next ranked server (whoami + 1); the `targets` array is a globally defined table that the balancer uses to do the migrations. Alternatively, to migrate all load a time step, we can remove the division by 2.

Figure **??** shows the performance of the modes using different migration units. Recall that this setup only has 2 sequencers and 2 servers, so performance may be different at scale. Even so, it is clear that client mode does not perform

as well for read-heavy workloads. We even see a throughput improvement when migrating all load off the first server, leaving the first server to do administrative tasks (this is common in the metadata cluster because the first server does a lot of the cache coherence work) while the second server does all the client work. Proxy mode does the best in both cases and shows large performance gains when completely decoupling client request handling and operation processing in Proxy Mode (Full). The parameter that controls the migration units helps the administrator control the sequencer co-location or distribution across the cluster. This trade-off was explored extensively in the Mantle paper but the experiments we present here are indicative of an even richer set of states to explore.

6.2.3 Feature: Backoff

Tuning the aggressiveness of the load balancer decision making is also a trade-off that administrators can control and explore. The balancing phase from 0 to 60 seconds in Figure 8 shows different degrees of aggressiveness in making migration decisions; CephFS makes a decision quickly and throughput jumps to 2500 ops/second within 10 seconds while Mantle takes more time to stabilize. This conservative behavior is controlled by programming the balancer to (1) use different conditions for when to migrate and (2) using a threshold for sustained overload.

We control the conditions for when to migrate using `when()`, a callback in the Mantle API. For the Mantle curve in Figure 8 we program `when()` to wait for load on the receiving server to fall below a threshold. This makes the balancer more conservative because it takes 60 seconds for cache coherence messages to settle. The Mantle curve in Figure 8 also takes longer to reach peak throughput because we want the policy to wait to see how migrations affect the system before proceeding; the balancer does a migration right before 50 seconds, realizes that there is a third underloaded server, and does another migration.

The other way to change aggressiveness of the decision making is to program into the balancer a threshold for sustained overload. This forces the balancer to wait a certain number of iterations after a migration before proceeding. In Mantle, the policy would use the save state function to do a countdown after a migation. Behavior graphs and performance numbers for this backoff feature is omitted for space considerations, but our experiments confirm that the more conservative the approach the less overall throughput.

Malacology pulls the load balancing out of the storage system to balance sequencers across a cluster. Compared to putting all sequencers on the same server, load balancing a sequencer using the proxy mode shows a $1.8\times$ throughput improvement. This latent capability also gives future programmers the ability to explore the different load balancing trade-offs including: load balancing modes to control forwarding vs. client redirection, load migration units to con-

trol sequencer distribution vs. co-location, and backoffs to control conservative vs. aggressive decision making.

7. Related Work

Programmability of operating systems and networking resources, including distributed storage systems is not new, but we are not aware of work that makes generalization of existing services into programmable resourses a key principle in storage systems design.

Programmable storage systems can be viewed as an infrastructure for creating abstractions to better separate policies from mechanisms. This idea is not new. Software-defined networks (SDNs) create such an abstraction by separating the control plane from the data plane (see for example [25]). This notion of control/data separation was also applied in software-defined storage (SDS) [33, 35]. Similarly, IOStack [19] is providing policy-based provisioning and filtering in OpenStack Swift. According to a SNIA white paper [14], the primary goal of SDS is to control and facilitate flexible and dynamic provisioning of storage resources of different kinds, including flash memory and disk drives, to create a virtualized mapping between common storage abstractions (e.g. files, objects, and blocks) and storage devices taking data service objectives in terms of protection, availability, performance, and security into account. A programmable storage system exposes internal abstractions so that end users (not necessarily operators) can create new services on top of the storage stack. Thus, our notion of programmable storage differs from "software-defined storage" (SDS) in terms of goals and scope, although definitions of SDS are still in flux.

Another view of programmable storage systems is one of tailoring systems resources to applications [4]. Related work includes work around the Exokernel [16] and SPIN [10] and Vino [29], the latter two addressed the ability of injecting code into the kernel to specialize resource management. Another approach is to pass hints between the different layers of the IO stack to bridge the semantic gap between applications and storage [4, 27, 31]).

Malacology uses the same Active and Typed Storage module presented in DataMods [38]; Asynchronous Service and File Manifolds can be implemented with small changes to the Malacology framework, namely asynchronous object calls and Lua stubs in the inode, respectively.

8. Conclusion and Future Work

Programmable storage is a viable method for eliminating duplication of complex error prone software that are used as workarounds for storage system deficiencies. However, this duplication has real-world problems related to reliability. We propose that systems expose their services in a safe way allowing application developers to customize system behavior to meet their needs while not sacrificing correctness. To illustrate the benefits of this approach we presented Mala-

cology, a programmable storage system that facilitates the construction of new services by re-purposing existing subsystem abstractions of the storage stack.

Future work will focus on constructing more interfaces to support a wide variety of storage system services that can be configured on-the-fly in existing systems. This work is one point along that path to producing general-purpose storage systems that can target special-purpose applications. Ultimately we want to utilize declarative methods for expressing new services.

References

[1] Zerovm and openstack swift.

[2] Apache Parquet Contributors. Parquet Columnar Storage Format, 2014.

[3] Michael Armbrust, Armando Fox, Rean Griffith, Anthony D. Joseph, Randy Katz, Andy Konwinski, Gunho Lee, David Patterson, Ariel Rabkin, Ion Stoica, and Matei Zaharia. A view of cloud computing. 53, 2010.

[4] A. C. Arpaci-Dusseau and R. H. Arpaci-Dusseau. Information and control in gray-box systems. In *SOSP '01*, Banff, Alberta, Canada, 2001.

[5] Mahesh Balakrishnan et al. Tango: Distributed data structures over a shared log. In *SOSP '13*, Farmington, PA, November 3-6 2013.

[6] Mahesh Balakrishnan, Dahlia Malkhi, Vijayan Prabhakaran, Ted Wobber, Michael Wei, and John D. Davis. Corfu: A shared log design for flash clusters. In *NSDI'12*, San Jose, CA, April 2012.

[7] Philip A. Bernstein, Sudipto Das, Bailu Ding, and Markus Pilman. Optimizing optimistic concurrency control for tree-structured, log-structured databases. In *SIGMOD '15*, 2015.

[8] Philip A. Bernstein, Colin W. Reid, and Sudipto Das. Hyder – a transactional record manager for shared flash. In *CIDR '11*, Asilomar, CA, January 9-12 2011.

[9] Philip A. Bernstein, Colin W. Reid, Ming Wu, and Xinhao Yuan. Optimistic concurrency control by melding trees. In *VLDB '11*, 2011.

[10] B. N. Bershad et al. Extensibility safety and performance in the spin operating system. In *SOSP '95*, Copper Mountain, CO, 1995.

[11] William J. Bolosky.

[12] Eric Brewer, Lawrence Ying, Lawrence Greenfield, Robert Cypher, and Theodore T'so. Disks for Data Centers. Technical Report, Google, 2016.

[13] Mike Burrows. The Chubby Lock Service for Loosely-Coupled Distributed Systems. In *Proceedings of the 7th symposium on Operating systems design and implementation*, OSDI '06.

[14] Mark Carlson, Alan Yoder, Leah Schoeb, Don Deel, Carlos Pratt, Chris Lionetti, and Doug Voigt. Software defined storage.

[15] Ceph Contributors. Ceph Documentation, 2010.

[16] D. R. Engler, M. F. Kaashoek, and Jr. J. O'Toole. Exokernel: An operating system architecture for application-level resource management. In *SOSP '95*, Copper Mountain, CO, 1995.

[17] Epping, Duncan and Denneman, Frank. VMware vSphere 5.1 Clustering Deepdive, accessed 03/21/2014, http://www.vmware.com/product/drs.

[18] Roxana Geambasu, Amit A. Levy, Tadayoshi Kohno, Arvind Krishnamurthy, and Henry M. Levy. Comet: an active distributed key-value store. In *Proceedings of the 9th USENIX conference on Operating systems design and implementation*, OSDI'10.

[19] Raúl Gracia-Tinedo et al. Iostack: Software-de ned object storage. *IEEE Internet Computing*, 20(3):10–18, May-June 2016.

[20] Matthias Grawinkel, Tim Sub, Gregor Best, Ivan Popov, and Andre Brinkmann. Towards Dynamic Scripted pNFS Layouts. In *Proceedings of the 2012 SC Companion: High Performance Computing, Networking Storage and Analysis*, SCC '12, 2012.

[21] Jim Gray. Tape is dead, disk is tape, flash is disk, RAM locality is king. *CIDR 2007 - Gong Show Presentation*, 2007.

[22] Jim Gray and Bob Fitzgerald. Flash Disk Opportunity for Server Applications. 6, 2008.

[23] Ajay Gulati, Ganesha Shanmuganathan, Anne Holler, and Irfan Ahmad. Cloud-scale resource management: challenges and techniques. In *Proceedings of the 3rd USENIX conference on Hot topics in cloud computing*, HotCloud'11, pages 3–3, Berkeley, CA, USA, 2011. USENIX Association.

[24] Patrick Hunt, Mahadev Konar, Flavio P. Junqueira, and Benjamin Reed. ZooKeeper: wait-free coordination for internet-scale systems. In *Proceedings of the 2010 USENIX conference on USENIX annual technical conference*, USENIXATC'10.

[25] S. Jain et al. B4: Experience with a globally-deployed software defined wan. In *SIGCOMM '13*, Hong Kong, China, 2013.

[26] Linux Foundation. Kinetic Open Storage Project, 2015.

[27] Michael Mesnier, Feng Chen, and Jason B. Akers. Differentiated storage services. In *SOSP '11*, Cascais, Portugal, October 23-26 2011.

[28] Erik Riedel, Garth A. Gibson, and Christos Faloutsos. Active storage for large-scale data mining and multimedia. In *24th international Conference on Very Large Databases (VLDB '98)*, New York, NY, 1998.

[29] M. I. Seltzer, Y. Endo, C. Small, and K. A. Smith. Dealing with disaster: Surviving misbehaved kernel extensions. In *OSDI'96*, Seattle, WA, 1996.

[30] Michael A. Sevilla, Noah Watkins, Carlos Maltzahn, Ike Nassi, Scott A. Brandt, Sage A. Weil, Greg Farnum, and Sam Fineberg. Mantle: A programmable metadata load balancer for the ceph file system. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, SC '15, 2015.

[31] Muthian Sivathanu, Vijayan Prabhakaran, Florentina I. Popovici, Timothy E. Denehy, Andrea C. Arpaci-Dusseau, and Remzi H. Arpaci-Dussea. Semantically-smart disk systems. San Francisco, CA, March 2003.

[32] SNIA. Implementing Multiple Cloud Storage APIs, November 2014.

[33] Ioan Stefanovici, Bianca Schroeder, Greg O'Shea, and Eno Thereska. sroute: Treating the storage stack like a network. In *FAST '16*, Santa Clara, CA, 2016.

[34] Ceph Team. Ceph's new monitor changes.

[35] Eno Thereska, Hitesh Ballani, Greg O'Shea, Thomas Karagiannis, Antony Rowstron, Tom Talpey, Richard Black, and Timothy Zhu. Ioflow: A software-defined storage architecture. In *SOSP '13*, Farmington, PA, November 3-6 2013.

[36] Lourival Vieira Neto, Roberto Ierusalimschy, Ana Lúcia de Moura, and Marc Balmer. Scriptable Operating Systems with Lua. In *Proceedings of the 10th ACM Symposium on Dynamic Languages*, DLS '14, pages 2–10, New York, NY, USA, 2014. ACM.

[37] Noah Watkins.

[38] Noah Watkins, Carlos Maltzahn, Scott Brandt, and Adam Manzanares. DataMods: Programmable File System Services. In *PDSW'12*, 2012.

[39] Noah Watkins, Carlos Maltzahn, Scott Brandt, Ian Pye, and Adam Manzanares. In-Vivo Storage System Development. In *Euro-Par 2013: Parallel Processing Workshops*, pages 23–32. Springer, 2013.

[40] Sage A. Weil, Andrew W. Leung, Scott A. Brandt, and Carlos Maltzahn. RADOS: a scalable, reliable storage service for petabyte-scale storage clusters. In *Proceedings of the 2nd international workshop on Petascale data storage: held in conjunction with Supercomputing '07*, PDSW '07, 2007.

[41] Sage A. Weil, Kristal T. Pollack, Scott A. Brandt, and Ethan L. Miller. Dynamic Metadata Management for Petabyte-Scale File Systems. In *Proceedings of the 17th ACM/IEEE Conference on Supercomputing*, SC'04, 2004.