

Malacology: A Programmable Storage System

Paper # 123

Abstract

Storage systems are caught between rapidly changing data processing systems and the increasing speed of storage devices. This puts tremendous pressure on storage systems to adapt both in terms of their interfaces and their performance. But adapting storage systems can be difficult because unprincipled changes might jeopardize years of code-hardening and performance optimization efforts that were necessary for users to entrust their data to the storage system. We introduce Malacology, a prototype programmable storage system to explore how existing abstractions of common services found in storage systems can be leveraged to address new data processing systems and the increasing speed of storage devices. This approach allows unprecedented flexibility for storage systems to evolve without sacrificing the robustness of its code-hardened subsystems. We illustrate the advantages and challenges of programmability by constructing two services out of existing abstractions: a file system metadata load balancer and a high-performance distributed shared-log that leverages flash devices. The evaluation demonstrates that our services inherit desirable qualities of the back-end storage system, including the ability to make trade-offs for lease capabilities, balance load (2× improvement), propagate interfaces (average of 222ms) and recover (15 second down time).

1 Introduction

A storage system implements abstractions designed to persistently store data and must exhibit a high level of correctness to prevent data loss. Storage systems have evolved around storage devices that often were orders of magnitude slower than CPU and memory, and therefore could dominate overall performance if not used carefully. Over the last few decades members of the storage systems community have developed ingenious strategies to meet correctness requirements while somewhat hiding the latency of traditional storage media [7]. To avoid lock-in by a particular vendor, users of storage systems have preferred systems with highly standardized APIs and lowest common denominator abstract data types such as blocks

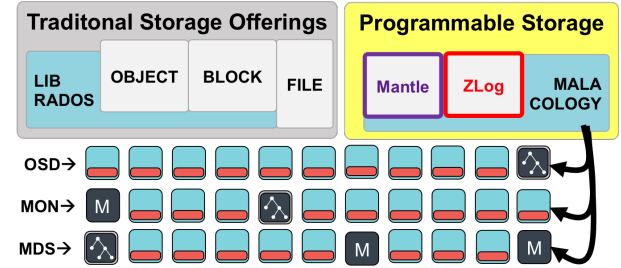


Figure 1: Scalable storage systems consist of object storage daemons (OSD) which store data fragments, monitor daemons (MON) that maintain cluster state, and service-specific daemons such as file system metadata servers (MDS). Malacology enables the programmability of internal abstractions and services (indicated by bold arrows) in order to leverage the back-end system’s durability, consistent versioning, consensus, and metadata subsystems. With Malacology, we add two new services, ZLog and Mantle, that sit alongside the traditional user-facing APIs (file, block, and object) in the storage stack.

of bytes and byte stream files [2].

A number of recent developments are disrupting traditional storage systems: (1) the falling prices of flash storage and the availability of new types of non-volatile memory that are orders of magnitude faster than traditional spinning media are moving overall performance bottlenecks away from storage devices to CPUs and networking, and pressure storage systems to shorten their code paths and incorporate new optimizations [15, 16]; (2) demand for managing structured data and flexible consistency semantics at scale pressure big data processing systems to use storage abstractions that can meet these demands [1]; and (3) production-quality scalable storage systems available as open source software have established and are continuing to establish new, *de-facto* API standards at a faster pace than traditional standards bodies [23, 18].

These three trends put evolutionary pressure on storage systems and raise the question whether there are principles that storage systems designers can follow to evolve storage systems efficiently and without jeopardizing years of code-hardening and performance optimization efforts that are important for users to continue to entrust their data to the storage system.

In this paper we investigate an approach that focuses on

generalizing existing storage system resources, services, and abstractions that in generalized form can be used to *program* new services. By doing so one can reuse subsystems and leverage their optimizations, established correctness, robustness, and efficiency. We will refer to this programmability as *programmable storage*, which differs from *active storage* (the injection and execution of arbitrary code in a storage system) and *software-defined storage* (the control of thin-provisioning of storage).

To illustrate the benefits and challenges of this approach we examine the programmability of Ceph [30, 29], the increasingly popular, production-level open-source distributed storage system. Something of a storage swiss army knife, Ceph supports file, block, and object interfaces simultaneously in a single cluster [9]. Ceph’s Reliable Autonomous Distributed Object Storage (RADOS) system is a cluster of storage devices (OSDs) that provide Ceph with data durability and integrity using replication, erasure-coding, and scrubbing [31].

Ceph also provides empirical evidence that developers embrace programmable storage: Ceph already supports the implementation of domain-specific logical object interfaces by composing existing interfaces. These interfaces are written in C++ and are statically loaded into the system. Figure 2 shows a dramatic growth in the production use of domain-specific interfaces in the Ceph community since 2010. What is most remarkable is that this trend contradicts the notion that API changes are a burden for users. Rather it appears that a gap in existing interfaces are being addressed through programmability. In fact, Table 1 categorizes existing interfaces and we clearly see a trend towards reusable services.




Figure 2: [source] Since 2010, the growth in the number of co-designed object storage interfaces in Ceph has been accelerating. This plot is the number of object classes (a group of interfaces), and the total number of methods (the actual API end-points).

The popularity of the custom object interface facility of

Category	Specialization	#
Locking	Shared	6
	Exclusive	
Logging	Replica	3
	State	4
	Timestamped	4
Metadata Management	RADOS Block Device	37
	RADOS Gateway	27
	User	5
	Version	5
Garbage Collection	Reference Counting	4

Table 1: A variety of RADOS object storage classes exist to expose interfaces to applications. # is the number of methods that use these categories.

Ceph hint at three trends in the Ceph community: (1) increasingly, the default algorithms/tunables of the storage system are insufficient for the application’s performance goals, (2) programmers are becoming more aware of their application’s behavior, and (3) programmers know how to manage resources to improve performance. Programmers gravitate towards object interfaces because it gives them the ability to tell the storage system about their application: if it is CPU or IO bound, if it has locality, if its size has the potential to overload a single proxy node, etc. The programmers know what the problem is and how to solve it, but until the ability to modify object interfaces, they had no way to express to the storage system how to handle their data.

To illustrate the benefits and challenges of this approach we have designed and evaluated Malacology, a programmable storage system for constructing new services by re-purposing existing subsystem abstractions contained in Ceph. We build the framework in Ceph by leveraging a broad spectrum of existing services, including distributed locking and caching services provided by MDSs, durability and logical storage devices provided by RADOS, and propagation of consistent cluster state provided by the monitoring service (see Figure 1). As we will show in this paper, this framework is expressive enough to provide the functionality necessary for implementing new services. Our contributions are:

A programmable storage system that re-uses and extends existing abstractions and services of Ceph, a common production-level storage system. It includes:

1. An extension of file system metadata by a new type of “sequencer” inode that together with Ceph’s capability subsystem implements high-performance, non-durable access serialization:

2. The use of Ceph’s PAXOS-based consensus infrastructure for versioning of metadata service load-balancing policies
3. The use of storage objects for persisting policies
4. An extension of the object interface classes for new logical storage/metadata devices

A distributed shared log service to demonstrate the feasibility of services using the programmable storage approach:

1. A high-performance distributed shared log service called ZLog which is based on CORFU [5]
2. A re-implementation of the programmable Mantle metadata load balancing service [21] using Malacology
3. An application of Mantle to manage load-balancing of sequencers.

The remainder of this paper is structured as follows. First we describe Malacology by presenting the subsystems within Ceph that we re-purpose, and briefly describe how those system are used within Malacology (§3). Next we describe the services that we have constructed within the Malacology framework (§4), and evaluate our ideas within our prototype implementation (§5). Finally we discuss related work and conclude.

2 Challenges

Establishing the infrastructure for programmability into existing services and abstractions of distributed storage systems is challenging, even if one assumes that the source code of the storage system and the necessary expertise for understanding it is available:

- Storage systems are generally required to be highly available so that any complete restarts of the storage system to reprogram them is usually unacceptable.
- Policies and optimizations are usually hard-wired into the services and one has to be careful when factoring them out not to introduce additional bugs.
- Mechanisms that are only exercised according to hard-wired policies and not in their full generality have hidden bugs that are revealed as soon as those mechanisms are governed by different policies. In our experience introducing programmability into a storage system proved to be a great debugging tool.

- Programmability, especially in live systems, implies changes that need to be carefully managed by the system itself, including versioning and propagation of those changes without affecting correctness.

3 Malacology

Malacology is both a prototype for a programmable storage system and a design approach to evolve storage systems efficiently and without jeopardizing years of code-hardening and performance optimization efforts. The guiding principle is to re-use existing services and extend them so that these services can be *programmed*. We accomplish programmability of a service by exporting bindings (or “hooks”) for an interpreted programming language so that programming can occur without having to restart the storage system (see also below, §3.4).

There are multiple reasons that make Lua an attractive runtime for implementing Malacology. Lua is a portable, embedded scripting language that offers superior performance and productivity trade-offs, including a JIT-based implementation that is well known for near native performance. Additionally, Lua has been used extensively in game engines, and systems research [26], including storage systems where it has been effectively used both on [14, 28, 12] and off [21] the performance critical path. Finally, the flexibility of the runtime allows to sandbox applications in order to address security and performance concerns.

We will now discuss the common subsystems used to manage storage system and how Malacology makes them programmable.

3.1 Cluster State Management

Keeping track of the state of a distributed system is an essential part of any successful service and a necessary component in order to diagnose and detect failures, when they occur. In the case of Ceph, a consistent view of cluster state among server daemons and clients is critical to provide strong consistency guarantees to clients. Ceph maintains cluster state information in per-subsystem “maps” (e.g. OSD map, MDS map) that record membership and status information. A PAXOS-based monitoring service (termed, the Monitor) is responsible for integrating state changes into cluster maps, responding to requests from out-of-date clients and synchronizing members of the cluster whenever there is a change in a map so that they all observe the same system state. As a fundamental building block of many system designs, consensus abstractions such as PAXOS are a common technique for

maintaining consistent data versions, and are a useful system to expose.

The default behavior of the monitor can be seen as a PAXOS-based notification system, similar to the one introduced in [8], allowing clients to identify when new values (termed epochs in Ceph) are associated to given maps. While Ceph does not expose this service directly, it does allow limited use by clients through a key-value service intended for low-volume, small size read-only configuration information. While a key-value service is generally useful, it doesn't provide many of the useful services hidden within the monitoring framework that would be required for applications with more demanding requirements, such as creating arbitrary maps or associating application-specific logic that can be executed for particular maps (or values in a map).

Malacology: Malacology exposes Ceph's PAXOS-based monitoring service implementing a generic API for adding arbitrary values to existing subsystem cluster maps, namely to the OSD map and the MDS map. As a consequence of this, applications can define simple but useful domain-specific logic to the PAXOS state machine, such as authorization control (just specific clients can write new values) or to trigger actions based on specific values (e.g. sanitize values). Since the monitor is intended to be out of high-performance I/O paths, a general guideline is to make use of this functionality infrequently and to assign small values to maps. Malacology itself makes use of this functionality to register, version and propagate dynamic code (Lua scripts) for new object interfaces defined in OSDs (§3.2) and load balancing policies in MDSs (§3.3).

3.2 Logical Storage Devices

Briefly described in Section 1, Ceph supports the creation of application-specific object interfaces [31]. The ability to offload computation can reduce data movement, and transnational interfaces can significantly simplify construction of complex storage interfaces that require uncoordinated parallel access.

An object interface is a plugin structured in a similar way to that of an RPC in which a developer creates a named function within the cluster that clients may invoke. In the case of Ceph each function is implicitly run within the context of an object specified when the function is invoked. Developers of object interfaces express behavior by creating a composition of native interfaces or other custom object interfaces, and handle serialization of function input and output. A wide range of native interfaces

are available to developers, such as reading and writing to a byte stream, controlling snapshots and clones, all of which may be transactionally composed with access to a sorted key-value service allowing applications to create rich interfaces, such as using a read-modify-write guard to protect a shared index defined over content store in a byte stream.




Figure 3: Malacology re-uses Ceph subsystems to allow users to dynamically define new object interfaces and load balancing policies. It achieves this by extending the OSD and MDS subsystems with an embedded Lua VM. To manage newly dynamic interface classes, it exposes the core functionality of the Monitor to allow arbitrary data to be associated to OSD and MDS cluster maps. After new interfaces are submitted to the monitor and the associated Paxos proposal is accepted by the monitor quorum, the script is redundantly and consistently distributed to OSDs so this snippets persist.

Malacology: The implementation of Ceph's object abstraction, although powerful, has some drawbacks. Supporting only C/C++ for object interface developers, Ceph requires distribution of compiled binaries for the correct architecture, adding a large barrier of entry for developers and system administrators. Second, having no way to dynamically unload modules, any changes require a full restart of an OSD which may have serious performance impacts due to a loss of cached data. And finally, the security limitations of the framework limit the use of object classes to all but those with administrative level access and deep technical expertise.

To address these concerns, Malacology takes advantage of community contributed Lua extensions to Ceph allowing new object interfaces to be dynamically loaded into the system and modified at runtime, resulting in a object storage API with economy of expression, which at the same time provides the full set of features of the base object class. New object interfaces that are expressed in thousands of lines of code can be implemented in approximately an order of magnitude less code [12].

Interface management. All of these modifications work in tandem to implement the desired behavior, and

are typically co-designed together such that each depends on the other to behave as expected. In highly dynamic environments such as a distributed storage system, changes to the system propagate at variable rates, requiring all components of the system to be able to adapt. Thus, system services realized through extensibility should be versioned together as a group, exposing the version to allow introspection and domain-specific treatment of version mismatch scenarios.

And finally, it is important to not underestimate the importance of the preservation of the modifications that enable a new system service. In many (if not all) cases, interfaces defining access to data are just as important as the data itself by virtue of inherently providing structural context. Thus, all components of a system service must be kept to the same standards of protection as data in the system itself.

To this end we have extended the monitoring service using the framework described in Section 3.1 to provide an interface for registration and verification of dynamically defined object interfaces. Using this service guarantees that interface definitions are not only made durable, but are transparently and consistently propagated throughout the cluster and that clients are properly synchronized with the latest interfaces.

3.3 Distributed Metadata Services

The distributed metadata service in Ceph provides clients with a POSIX file system abstraction [32]. In general, distributed file systems protect resources by providing hierarchical indexing and distributed locking services. In Ceph, the locking service implements a capability-based system that express what data and metadata clients are allowed to access as well as what state they may cache and modify locally. While designed for a fixed file abstraction, indexing, locking, and caching are all common services that are useful to a broad spectrum of applications.

Ceph also addresses the challenge of balancing metadata load using load balancing policies that compute metrics based on system state (e.g. CPU and memory utilization) and statistics collected by the cluster (e.g. the popularity of an inode). Ceph uses dynamic subtree partitioning to move variable sized namespace subtrees. These units can be shipped anywhere (i.e., to any MDS of any capacity) at any time for any reason. The original balancer was designed with hard-coded policies and tunables. Distributed applications that share centralized resources (e.g. a database or directory) face similar challenges which are often solved using application-specific sharding.

Malacology: New interface hooks are added in strategic spots in the metadata service to control load balancing and capabilities, and for defining new file types allowing the indexing service to be re-purposed for other types of named resources.

Load balancing. Mantle [21], a programmable metadata load balancer, has been re-implemented on Malacology so that its infrastructure can enjoy all the properties of other Ceph components. The load balancing logic has stayed largely the same but the infrastructure has been improved to safely load, manage, and persist its load balancing policies.

Safely Loading Classes. The original Mantle API is fragile: the API is not enforced by the runtime, the system allows strings to be injected through the admin daemon, constructing the Lua script in C++ is clunky, and the administrator can inject really bad policies (e.g., while 1) that brings the whole system down. Malacology exposes the important subsystems of Ceph’s logical storage device framework so that the MDS can safely and dynamically load policies specified as Lua classes.

Policy management. A new interface is added to the metadata service subsystem using the monitoring framework to support management of interfaces. This ensures that the correct balancer version is consistently distributed and verified on all MDS nodes. Policies are also durably saved with Ceph’s replication and redundancy.

Capabilities. The metadata service also manages sessions with its clients. This gives both subsystems (MDSs and clients) the ability to cache data and grab locks. The metadata service is responsible for responding to client requests for capabilities by first revoking caps from other clients. Currently, clients voluntarily release their capabilities, and the MDS maintains a queue of requests. Both clients and MDSs use a best-effort approach to revoking or releasing capabilities.

The clients retain the best-effort, voluntary release of capabilities, but in Malacology, the queuing policies are generalized within the MDS to allow for centralized policies such as fairness or priority.

File Types. We leverage the same active and typed storage module presented in DataMods [27] to assign state and a set of functions to a file (stored in the inode). This is valuable for metadata load balancing so that certain types of files can have a higher priority when making balancing decisions.

3.4 Durability

Ceph provides storage by striping and replicating data across RADOS [31], the reliable distributed object store.

Service	Ceph Component	Used for...
Mantle	RADOS MON maps	policy durability policy versioning
ZLog	Logical Storage MON maps	CORFU protocol interface install
Seq.	MDS file types MDS fault toler.	log management seq. fault toler.

Table 2: To implement the Mantle and ZLog services, we leverage Ceph components. Malacology uses these subsystems because they provide powerful abstractions.

RADOS uses many techniques to ensure that data is not corrupted or lost, such as erasure coding, replication, and data scrubbing. Furthermore, many of these techniques try to be autonomous so that work is distributed across the cluster. For example, when placement groups change, the OSDs re-balance and re-shard data in the background in a process called placement group splitting.

In order to reduce load on the monitoring service, Ceph OSDs use a gossip protocol to efficiently propagate changes to cluster maps throughout the system, and autonomously initiate recovery mechanisms when failures are discovered.

Malacology: Metadata service policies and object storage interfaces are stored durability within RADOS and managed by storing references with the object maps. Since the cluster already propagates a consistent view of these data structures, we use this service to automatically install interfaces in OSDs, and install policies within the MDS such that clients and daemons are synchronized on correct implementations without restarting.

4 Services Built on Malacology

In this section we describe two services built on top of Malacology. The first is Mantle, a framework for dynamically specifying metadata load balancing policies. The second system, ZLog, is a high-performance distributed shared-log. In addition to these services, we’ll demonstrate how we combine ZLog and Mantle to implement service-aware metadata load balancing policies.

4.1 Mantle: Programmable Load Balancer

Mantle is a programmable load balancer that separates the metadata balancing policies from their mechanisms. Administrators inject code to change how the metadata cluster distributes metadata. In [21] the authors showed how using Mantle, one can implement a single node metadata

service, a distributed metadata services with hashing, and a distributed metadata service with dynamic subtree partitioning.

The original implementation was “hard-coded” into Ceph and lacked robustness (lacking things such as versioning, durability, or reliable policy distribution). Re-implemented using Malacology, Mantle now enjoys (1) the versioning of MONs and (2) the durability and distribution provided by RADOS. The version of the load balancer, which is stored in the MDS map, corresponds to an object name in RADOS that houses the load balancer code.

4.1.1 Versioning Balancer Policies

Ensuring that the version of the current load balancer is consistent across the MDS cluster was ignored in the original implementation. The user had to set the version on each individual MDS and it was trivial to make the version inconsistent across MDS nodes.

With Malacology, Mantle stores the version of the current load balancer in the MDS map. Similar to the PG map or MON maps, the MDS map has all the information about the topology and state of the MDS cluster. Recall that the MONs propagate these maps across the cluster and ensure that the map versions are consistent with PAXOS. While storing the code snippets that define the policies in the MDS map is possible, it is not recommended as it is an expensive operation (a PAXOS voting round).

The user changes the version of the load balancer using a new CLI command:

```
fs set lua.balancer.class <version>
```

4.1.2 Making Balancer Policies Durable

The load balancer version described above corresponds to the name of an object in RADOS that holds the actual Lua balancing code. The objects containing load balancer code are stored in a system-maintained group of objects (i.e. a pool). When MDS nodes start balancing load, they first check the latest version from the MDS map and compare it to the balancer they have loaded. If the version has changed, they dereference the pointer to the balancer version by reading the corresponding object in RADOS. This is in contrast to the original balancer which stored load balancer code on the local file system – a technique which is unreliable and results in inconsistencies if any MDS accidentally modifies the code.

The balancer pulls the Lua code from RADOS synchronously; asynchronous reads are not possible because

of the architecture of the MDS. The synchronous behavior is not the default behavior for RADOS operations, so we achieve this with a timeout: if the asynchronous read does not come back within half the balancing tick interval the operation is cancelled and a Connection Timeout error is returned. By default, the balancing tick interval is 10 seconds, so Mantle will use a 5 second second timeout.

This design allows Mantle to immediately return an error if anything RADOS-related goes wrong. We use this implementation because we do not want to do a blocking OSD read from inside the global MDS lock. Doing so would bring down the MDS cluster if any of the OSDs are not responsive.

Storing the balancers in RADOS is only possible because we use Lua as the language for writing balancer code. If we used a language that needs to be compiled, like the C++ object classes in the OSD, we would need to ensure binary compatibility, which is complicated by different operating systems, distributions, and compilers.

4.1.3 Logging, Debugging, and Warnings

In the original implementation, Mantle would log all errors, warnings, and debug messages to the local log on each MDS. To get the simplest status messages or to debug problems, the user would have to log into each MDS individually, look at the logs, and reason about their causality.

With Malacology, Mantle re-uses the centralized logging features of the MON subsystem. Important errors, warnings, and info messages are sent with the MON beacon subsystem and appear in the monitor cluster log so instead of users going to each node, they can watch messages appear at the monitor. Messages are used discretely, so as not to spam the monitor with frivolous debugging but the large events, like balancer version changes or failed subsystems show up in the centralized log.

4.2 ZLog: A Fast Distributed Shared-Log

The second service implemented on Malacology is ZLog, a high-performance distributed shared-log that is based on the CORFU protocol [5]. A shared-log is a general, yet powerful abstraction used to build many distributed systems. However, existing implementations that rely on consensus algorithms such as PAXOS funnel I/O through a single point introducing a bottleneck that restricts throughput. In contrast, the CORFU protocol is able to achieve high throughput using a network counter service, called a *sequencer*, that avoids all persistent I/O in the common case.

While a full description of the CORFU system is beyond the scope of this paper, we'll briefly describe the custom storage device interface, sequencer service, and recovery protocol, and how these services are instantiated in the Malacology framework.

4.2.1 Sequencer

High-performance in CORFU is achieved using a sequencer service that assigns log positions to clients by reading from a volatile, in-memory counter which can run at a very high throughput and at low-latency. And since the sequencer is centralized, ensuring serializability in the common case is trivial. The primary challenge in CORFU is handling the failure of the sequencer in a way that preserves correctness. Failure of the sequencer service in CORFU is handled by a recovery algorithm that recomputes the new sequencer state using an application-specific custom storage interface to discover the tail of the log, while simultaneously invalidating stale client requests using an epoch-based protocol.

We implement the sequencer service in Malacology as an application-specific file type that associates a small amount of metadata, including the 64-bit log tail position, to the inode state managed by the Ceph metadata service. Since a sequencer is associated with a unique log, this approach has the added benefit of allowing the metadata service to also handle naming, by representing each log service instance within the standard POSIX hierarchical namespace.

Sequencer interface. The sequencer resource supports the ability to *read()* the current tail value as well get the *next()* position in the log which also atomically increments the tail position. We implement this functionality as methods on ZLog inode resources. The primary challenge in mapping the sequencer resource to the metadata service is handling serialization correctly.

Initially we sought to directly model the sequencer service in Ceph as a non-exclusive, non-cacheable resource, forcing clients to perform a round-trip access to the resource at the authoritative metadata server for the sequencer inode. Interestingly, we found that the capability system in Ceph uses a strategy to reduce metadata service load by allowing multiple clients that open a shared file to temporarily cache the inode resource, resulting in a round-robin, best-effort batching behavior. When a single client is accessing the sequencer resource it is able to cache state locally without interruption.

While unexpected, this discovery allowed us to explore an implementation strategy that we had not previously considered. In particular, for bursty clients, and clients that can tolerate increased latency, this mode of operation

may allow a system to achieve much higher throughput than a system with a centralized sequencer service. We utilize the programmability of the metadata service to define a new policy for handling capabilities that controls the amount of time that clients are able to cache the sequencer resource. This allows an administrator or application to control the trade-off between latency and throughput beyond the standard best-effort policy that is present in Ceph by default.

In Section 5 we quantify the trade-offs of throughput and latency for an approach based on a round-robin batching mode, and compare this mode to one in which the metadata server mediates access to the sequencer state when in it is being shared among multiple clients.

Balancing policies. As opposed to the batching mode for controlling access to the sequencer resource, more predictable latency can be achieved by treating the sequencer inode as a shared non-cacheable resource, forcing clients to make a round-trip to the metadata service. However, the shared nature of the metadata service may prevent the sequencer from achieving maximum throughput. To provide a solution to this issue we have taken advantage of the programmability of the metadata service load balancing to construct a service-specific load balancing policy. As opposed to a generic balancing policy that may strive for uniform load distribution, a ZLog-specific policy may utilize knowledge of inode types to migrate the sequencer service to provisioned hardware during periods of contention or high demand.

4.2.2 Storage Interface

The storage interface is a critical component in the CORFU design. Clients independently map log positions that they have obtained from a sequencer service (described in detail in the next section) onto storage devices, while storage devices provide an intelligent *write-once*, random *read* interface for accessing log entries. The key to correctness in CORFU lies with the enforcement of up-to-date epoch tags on client requests; requests tagged with out-of-date epoch values are rejected, and clients are expected to request a new tail from the sequencer after refreshing state from an auxiliary service. This mechanism forms the basis for sequencer recovery.

In order to repopulate the cached tail value during recovery of a sequencer, the maximum position in the log must be obtained. To do this, the storage interface exposes an additional *seal* method that atomically installs a new epoch value and returns the maximum log position that has been written. Since the recovery protocol does not inform clients of a new epoch value until the end of the process, the validity of the tail can be guaranteed de-

spite the iterative process of contacting storage devices. Recovery of the sequencer process itself may be handled in many ways, such as leader election using an auxiliary service such as PAXOS.

5 Evaluation

We have conducted our evaluation on CloudLab as well as our own 20 node cluster. We follow the “Popper Convention” [4] so as to try and make our experiments as reproducible as possible. We have a GitHub repository with our source code, experimental environment, configuration parameters, and workloads but due to the double-blind policy, we have temporarily disabled the links to our graphs and experimental details.

Much of this work focuses on the programmability of Ceph and the services that can be layered on top of Malacology. As such, we focus on the performance of the components and subsystems exposed by the framework. Detailed system-wide performance measurements and analysis is forthcoming for both Mantle and ZLog, as we try to find the best metadata balancing policies and physical design implementation trade-offs respectively.

5.1 ZLog

We evaluate the mapping of ZLog onto Malacology by examining the costs of utilizing the metadata service as a sequencer, and propagating interface changes throughout the cluster using the monitoring service. We have omitted evaluation of the overhead costs of adding object interfaces themselves due to lack of space, and the fact that they are heavily tested and used already in production with acceptable performance, introducing a small constant overhead per operation but otherwise does not affect scalability.

5.1.1 Sequencer

We evaluated the performance of the ZLog Sequencer implemented in several different configurations by showing a trade-off between aggregate throughput achieved (shown in Figure 4), and latency distribution (shown in Figure 5).

In Figure 4 the highest throughput labeled 1C shows that a single client caching the sequencer resource may obtain new sequences at a rate of nearly 6 million positions per second. This is not surprising given that all state is cached locally in the client. When this resource is shared between two clients, the MDS uses a best-effort approach to granting the sequencer resource to each client

(value $2C, BE$). This sharing has a large impact on performance, as shown by the nearly three orders of magnitude drop off in throughput.




Figure 4: [source] Sequencer throughput by re-using various services. The highest performance is achieved using a single client with exclusive, cacheable privilege. Round-robin sharing of the sequencer resource is affected by the amount of time the resource is held, with best-effort performing the worst.




Figure 5: [source] The latency distribution shows that if the clients hold their capabilities longer, performance decreases. Changing the delay via policy can help the sequencer and clients strike a balance between throughput and latency.

While a best-effort approach to sharing capabilities makes sense for a file system client, throughput drops significantly with two clients as the majority of the time neither client has the capability as it ping-pongs between clients. This observation reveals a tuning parameter: how long a client maintains its capability lease. In throughput labeled $2C, <time>$ we show aggregate throughput as we decrease the amount of time that a client holds its

lease before responding to a request for the capability. At lease times as low as 10ms throughput achieve a significant proportion of the locally cached throughput.

The Malacology Approach: Re-use the capability management features of the MDS. While the batching mechanism can achieve high-throughput, it is not appropriate for applications that require predictable latency. The throughput in Figure 4 labeled $2C, ZLog$ corresponds to a centralized sequencer implementation. While throughput is lower for two clients (because of the synchronous network round-trip required for each request), Figure 5 shows that this approach achieves more predictable latency compared to all other approaches except for the single client locally cached mode.

5.1.2 Interface Propagation

In Section 3.1 we described an interface management service that is responsible for the durable and consistent cluster-wide distribution of interface versions. We evaluate the costs of such a service by embedding interface definitions in a Ceph cluster membership data structure that is managed by a PAXOS instance, and distributed using a scalable, peer-to-peer gossip protocol that is built into Ceph. We assume an architecture that consists of a light-weight interface that enforces the semantics of interface management resulting in one or more interface updates that must be consistent across the cluster. The following experiments focus on the cost of achieving distribution of consistent a interface version.




Figure 6: [source] Cluster-wide interface update latency, excluding the Paxos proposal cost for interface commit. Each of the 120 OSDs used a RAM-based object store, while the 3 OSDs used an HDD-backed object store.

Figure 6 shows the CDF of the latency of installing and distributing an interface update on all OSDs in the cluster. The latency is defined as the elapsed time following the PAXOS proposal for an interface update until each OSD makes the update live (the cost of the PAXOS proposal is

configurable and is discussed below). The latency measurements were taken on the nodes running Ceph server daemons, and thus exclude the client round-trip cost. In each of the experiments 1000 interface updates were observed.

The first experiment shown in Figure 6 illustrates a lower bound cost for updates in a large cluster by avoiding disk I/O for storing interface updates locally within each OSD. In the experiment labeled "120 OSD (RAM)" a cluster of 120 OSDs (10 OSDs x 12 servers) using an in-memory data store were deployed, showing a latency of less than 54 ms with a probability of 90% and a worst case latency of 194 ms. The curve labeled "3 OSD (HDD)" is a smaller cluster where each of the three OSDs uses a spinning disk as the object store that saves a copy of the cluster map, and shows a latency of 37 ms with a 90% probability and a worst case latency of 255 ms. Since the cost of a disk I/O is larger than the same I/O executed against the in-memory device, the predominant cost in a lightly loaded cluster appears to be the number of OSDs.

The Malacology Approach: Re-use Ceph's Monitor functionality to propagate application-specific state that needs to be highly-available and strongly consistent. The real cost of interface management, in addition to cluster-wide propagation of interface updates, includes the network round-trip to the interface management service, the PAXOS commit protocol, and other factors such as system load. By default PAXOS proposals occur periodically with a 1 second interval in order to accumulate updates. In a minimum, realistic quorum of 3 monitors using HDD-based storage, we were able to decrease this interval to an average of 222 ms.

5.2 MDS as a ZLog Sequencer

The following experiments demonstrate the feasibility of using the Ceph MDS cluster as a sequencer. The goal of the workloads is to saturate the MDS cluster, so we design sequencers uses a more heavy-weight file system operation. The experiments use our internal cluster and there are 15 clients and between 1 and 3 MDS nodes.

5.2.1 Load Balancing

We use a balancer policy specifically designed for the sequencer. The balancer aggressively sheds half its load to its neighbor as soon as it can. It is based off the "Greedy Spill Even" approach in the Mantle [21] paper but is reproduced here for the reader's convenience.

Figure 7 gives a macro-level view of the performance of the Ceph components exposed by Malacology. Each curve shows the throughput from the perspective of one




Figure 7: [source] A ZLog sequencer that uses the CephFS MDS enjoys the load balancing capabilities of the metadata cluster. The "Sequencer Baseline" has no background sequencers, while the "2 Seq" curves are the throughput of 1 sequencer when an additional sequencer is running in the background generating load.

```

— Metadata server load
mdsload = MDSs[i]["all"]

— When policy
if MDSs[whoami]["load"] > .01 and
   MDSs[whoami+1]["load"] < .01 then

— Where policy
targets[whoami+1] = allmetaload / 2

```

Figure 8: The sequencer balancing policy that works best for bursty workloads.

sequencer (y axis) over time (x axis). The "Sequencer Baseline" curve is one sequencer serving 7 clients. The sequencer operates on an MDS cluster where balancing is turned off and there are no other workloads running. "No Load Balancing" adds another sequencer of equal load running in the background. The performance of the background sequencer is omitted. The last two policies, "Sequencer Balancing Policy" and "Conservative Balancing Policy", show the performance of the sequencer (with a background sequencer running) when using Mantle with a 3 MDS node cluster. These policies differ by the threshold at which they start balancing: "Sequencer Balancing Policy" starts balancing immediately and "Conservative Balancing Policy" waits for 120 seconds.

Figure 9 gives a time series graph to show the behavior of the same workload, 2 sequencers, with balancing turned on and off. The curves are the throughput (y axis) over time (x axis) and the graph has been partitioned into




Figure 9: [source] The sequencer using Mantle in Figure 7 migrates to a new MDS immediately because it uses a policy designed for bursty workloads. This is a trace of the same workload running twice: once with a single MDS and once with Mantle.

the two cluster setups. The CPU utilization is confined to one core the MDS is not multi-threaded, as it has been shown that metadata performance does not improve with multiple threads [10].

Overall, the clients achieve lower performance than the capability experiments in Figure 4 because the sequencer uses a slower metadata operation to manage the capability. “Sequencer Baseline” is the best performance for this particular workload because there is no contention. Performance is stable and the MDS is under-loaded because we do not see the spike at the end like the other curves. It appears the bottleneck is somewhere with the clients – perhaps it is our 1Gbit/s network link. “2 Seq; No Load Balancing” shows an overloaded sequencer because the performance of our one sequencer is low and unstable (high variation in performance). We also deduce that it is overloaded by comparing to the “2 Seq; Mantle Load Balancing” curve, which shows improved performance when load is migrated. At 10 seconds, the sequencer balancer immediately starts balancing and decides that two sequencer streams should be split onto two MDSs. It migrates to its neighbor as per the policy.

The Malacology Approach: Re-use the Ceph metadata migration, load balancing, cache coherence, and capabilities subsystems of the MDS. Applying Mantle’s load balancing policy shows $2.02 \times$ performance improvement and 16% degradation in performance over the baseline, even though the baseline has half as many clients.

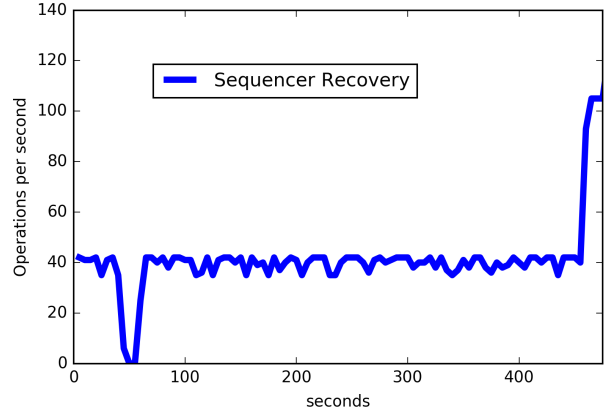


Figure 10: The sequencer using Mantle also inherits all the recovery protocols of the Ceph metadata cluster. The throughput of the sequencer drops when we terminal MDS 0. The performance quickly resumes and the sequencer is only down for 15 seconds.

5.2.2 Recovery

This experiment shows how a sequencer handing out tokens can use the Ceph MDS recovery protocols to recover. For this workload, two sequencers are running on a metadata cluster *without* multiple active MDS nodes. In this cluster, MDS 1 and MDS 2 are in “standby” mode, meaning one of them will step in should the head MDS 0 fail. At the 60 second mark, we manually kill MDS 0.

Figure 10 shows the sequencer throughput (y axis) over time (x axis). When MDS 0 dies, the MON cluster notices and quickly assigns a new MDS to handle the namespace. This is when MDS 1 takes over the sequencer duties and starts handing out tokens.

The Malacology Approach: Re-use the Ceph metadata protocols for failure detection and recovery. The sequencer recovers in 15 seconds on a new MDS and it maintains the same throughput that it had before the failure.

6 Related Work

Programmability of operating systems and networking resources, including distributed storage systems is not new, but we are not aware of work that makes generalization of existing services into programmable resources a key principle in storage systems design.

Programmable storage systems can be viewed as an infrastructure for creating abstractions to better separate policies from mechanisms. This idea is not new. Software-defined networks (SDNs) create such an abstraction by separating the control plane from the data

plane (see for example [17]). This notion of control/data separation was also applied in software-defined storage (SDS) [3, 25, 24]. Similarly, IOSTack [13] is providing policy-based provisioning and filtering in OpenStack Swift.

Another view of programmable storage systems is one of tailoring systems resources to applications [3]. Related work includes work around the Exokernel [11] and SPIN [6] and Vino [20], the latter two addressed the ability of injecting code into the kernel to specialize resource management. Another approach is to pass hints between the different layers of the IO stack to bridge the semantic gap between applications and storage [3, 22, 19]).

Malacology uses the same Active and Typed Storage module presented in DataMods [27]; Asynchronous Service and File Manifolds can be implemented with small changes to the Malacology framework, namely asynchronous object calls and Lua stubs in the inode, respectively.

7 Conclusion and Future Work

Programmable storage is a viable method for eliminating duplication of complex error prone software that are used as workarounds for storage system deficiencies. However, this duplication has real-world problems related to reliability. We propose that system expose their services in a safe way allowing application developers to customize system behavior to meet their needs while not sacrificing correctness.

We are intending to pursue this work towards the goal of constructing a set of customization points that allow a wide variety of storage system services to be configured on-the-fly in existing systems. This work is one point along that path in which we have looked at targeting special-purpose storage systems. Ultimately we want to utilize declarative methods for expressing new services.

References

- [1] Apache Parquet Contributors. Parquet Columnar Storage Format, 2014.
- [2] Michael Armbrust, Armando Fox, Rean Griffith, Anthony D. Joseph, Randy Katz, Andy Konwinski, Gunho Lee, David Patterson, Ariel Rabkin, Ion Stoica, and Matei Zaharia. A view of cloud computing. 53, 2010.
- [3] A. C. Arpaci-Dusseau and R. H. Arpaci-Dusseau. Information and control in gray-box systems. In *SOSP '01*, Banff, Alberta, Canada, 2001.
- [4] Anonymous Author. Omitted to blind submission. *Anonymous*, 1900.
- [5] Mahesh Balakrishnan, Dahlia Malkhi, Vijayan Prabhakaran, Ted Wobber, Michael Wei, and John D. Davis. Corfu: A shared log design for flash clusters. In *NSDI'12*, San Jose, CA, April 2012.
- [6] B. N. Bershad et al. Extensibility safety and performance in the spin operating system. In *SOSP '95*, Copper Mountain, CO, 1995.
- [7] Eric Brewer, Lawrence Ying, Lawrence Greenfield, Robert Cypher, and Theodore T'so. Disks for Data Centers. Technical Report, Google, 2016.
- [8] Mike Burrows. The Chubby Lock Service for Loosely-Coupled Distributed Systems. In *Proceedings of the 7th symposium on Operating systems design and implementation*, OSDI '06.
- [9] Ceph Contributors. Ceph Documentation, 2010.
- [10] Konstantinos Chasapis, Manuel F. Dolz, Michael Kuhn, and Thomas Ludwig. Evaluating Lustre's Metadata Server on a Multi-socket Platform. In *Proceedings of the 9th Parallel Data Storage Workshop*, PDSW '14, 2014.
- [11] D. R. Engler, M. F. Kaashoek, and Jr. J. O'Toole. Exokernel: An operating system architecture for application-level resource management. In *SOSP '95*, Copper Mountain, CO, 1995.
- [12] Roxana Geambasu, Amit A. Levy, Tadayoshi Kohno, Arvind Krishnamurthy, and Henry M. Levy. Comet: an active distributed key-value store. In *Proceedings of the 9th USENIX conference on Operating systems design and implementation*, OSDI'10.
- [13] Raúl Gracia-Tinedo et al. Iostack: Software-defined object storage. *IEEE Internet Computing*, 20(3):10–18, May-June 2016.
- [14] Matthias Grawinkel, Tim Sub, Gregor Best, Ivan Popov, and Andre Brinkmann. Towards Dynamic Scripted pNFS Layouts. In *Proceedings of the 2012 SC Companion: High Performance Computing, Networking Storage and Analysis*, SCC '12, 2012.
- [15] Jim Gray. Tape is dead, disk is tape, flash is disk, RAM locality is king. *CIDR 2007 - Gong Show Presentation*, 2007.

- [16] Jim Gray and Bob Fitzgerald. Flash Disk Opportunity for Server Applications. 6, 2008.
- [17] S. Jain et al. B4: Experience with a globally-deployed software defined wan. In *SIGCOMM '13*, Hong Kong, China, 2013.
- [18] Linux Foundation. Kinetic Open Storage Project, 2015.
- [19] Michael Mesnier, Feng Chen, and Jason B. Akers. Differentiated storage services. In *SOSP '11*, Cascais, Portugal, October 23-26 2011.
- [20] M. I. Seltzer, Y. Endo, C. Small, and K. A. Smith. Dealing with disaster: Surviving misbehaved kernel extensions. In *OSDI'96*, Seattle, WA, 1996.
- [21] Michael A. Sevilla, Noah Watkins, Carlos Maltzahn, Ike Nassi, Scott A. Brandt, Sage A. Weil, Greg Farnum, and Sam Fineberg. Mantle: A programmable metadata load balancer for the ceph file system. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, SC '15, 2015.
- [22] Muthian Sivathanu, Vijayan Prabhakaran, Florentina I. Popovici, Timothy E. Denehy, Andrea C. Arpaci-Dusseau, and Remzi H. Arpaci-Dusseau. Semantically-smart disk systems. San Francisco, CA, March 2003.
- [23] SNIA. Implementing Multiple Cloud Storage APIs, November 2014.
- [24] Ioan Stefanovici, Bianca Schroeder, Greg O'Shea, and Eno Thereska. srout: Treating the storage stack like a network. In *FAST '16*, Santa Clara, CA, 2016.
- [25] Eno Thereska, Hitesh Ballani, Greg O'Shea, Thomas Karagiannis, Antony Rowstron, Tom Talpey, Richard Black, and Timothy Zhu. Ioflow: A software-defined storage architecture. In *SOSP '13*, Farmington, PA, November 3-6 2013.
- [26] Lourival Vieira Neto, Roberto Ierusalimsky, Ana Lúcia de Moura, and Marc Balmer. Scriptable Operating Systems with Lua. In *Proceedings of the 10th ACM Symposium on Dynamic Languages*, DLS '14, pages 2–10, New York, NY, USA, 2014. ACM.
- [27] Noah Watkins, Carlos Maltzahn, Scott Brandt, and Adam Manzanares. DataMods: Programmable File System Services. In *PDSW'12*, 2012.
- [28] Noah Watkins, Carlos Maltzahn, Scott Brandt, Ian Pye, and Adam Manzanares. In-Vivo Storage System Development. In *Euro-Par 2013: Parallel Processing Workshops*, pages 23–32. Springer, 2013.
- [29] Sage Weil et al. Ceph.
- [30] Sage A. Weil, Scott A. Brandt, Ethan L. Miller, Darrell D. E. Long, and Carlos Maltzahn. Ceph: A Scalable, High-Performance Distributed File System. In *Proceedings of the 7th USENIX Symposium on Operating Systems Design & Implementation*, OSDI'06, 2006.
- [31] Sage A. Weil, Andrew W. Leung, Scott A. Brandt, and Carlos Maltzahn. RADOS: a scalable, reliable storage service for petabyte-scale storage clusters. In *Proceedings of the 2nd international workshop on Petascale data storage: held in conjunction with Supercomputing '07*, PDSW '07, 2007.
- [32] Sage A. Weil, Kristal T. Pollack, Scott A. Brandt, and Ethan L. Miller. Dynamic Metadata Management for Petabyte-Scale File Systems. In *Proceedings of the 17th ACM/IEEE Conference on Supercomputing*, SC'04, 2004.